

Peer to peer systems (P2P)

- Napster Structure
- client machines ("peers")
- servers - store the file name and the info about the file

Client

- connects to a Napster server
- search
 - sends message to the server
 - returns a list of host
 - client pings each host in the list to find the transfer rates
 - client fetches file from best host
- all communication uses TCP (Transmission Control Protocol)

Joining a P2P system

Problems

- centralized server a source of congestion
- centralized server single point of failure
- no security
- indirect infringement

Gnutella

- eliminate the servers
- client machines search and retrieve amongst themselves
- peers store their own files
- peers connected in an overlay graph
- routes messages within the overlay graph (types of payloads)
 - query (search)
 - queryHit (response to query)
 - ping (to probe network for other peers)
 - pong (reply to ping, contains address of another peer)
 - push (used to initiate file transfer)
- how to search for a file
 - descriptor id bytes - 0 to 15
 - id of this search transaction
 - payload descriptor - 15 to 16
 - type of payload
 - TTL (time to live) - 17
 - decremented at each hop, message is dropped when TTL = 0.
 - Hops - 18
 - incremented every time a message is forwarded
 - Payload length - 22
 - number of bytes of message following this header
- Query
 - contains the minimum speed
 - contains the search criteria

- QueryHit: successful result to a query
 - number of hits
 - port
 - information about the sender
 - ip_address
 - information about the sender
 - speed
 - information about the sender
 - (fileindex,filename,fsize)
 - information about the sender
 - results
 - servent_id
 - unique identifier of responder; a function of its IP address
- Avoiding excessive traffic
 - to avoid duplicate transmissions, each peer maintains a list of recently received messages
 - query forwarded to all neighbors except peer from which received
 - each query forwarded only once
 - QueryHit routed back only to peer from which Query received with same DescriptorID
 - for flooded messages, duplicates with same descriptorID and payload descriptor are dropped
 - QueryHit with DescriptorID for which Query not seen is dropped
- after receiving query hit messages
 - requestor choose best query hit responder
 - initiates HTTP request directly to responder's ip+port
 - responder then replies with file packet after this message
- HTTP file transfer protocol
 - because it's standard, well debugged, and widely used
- Why the range field in the GET request
 - to support partial file transfers
- what if responder is behind a firewall?
 - prevents messages from coming in
 - dealing with firewalls
 - requester sends push to responder asking for file transfer
 - push message
 - servent_id
 - same as in received query hit
 - fileindex
 - same as in received query hit
 - ip_address
 - address at which requestor can accept incoming connections
 - same as in received query hit
 - port
 - address at which requestor can accept incoming connections
 - responder establishes a TCP connection at ip_address, port specified
 - requestor then sends GET to responder (as before) and file is transferred as explained earlier
 - Ping-Pong
 - ping
 - no payload

- pong
 - port, ip_address, num files shared, number of KB shared
- peers initiate Ping's periodically
- Ping's flooded out like Query's, Pong's routed along reverse path like QueryHit's
- Pong replies used to update set of neighboring peers
 - to keep neighbor lists fresh in spite of peers joining, leaving and falling
- summary
 - peers/ servants maintain "neighbors", this forms an overlay graph
 - queries flooded out, ttl restricted
 - QueryHit (replies) reverse path routed
 - supports file transfer through firewalls
 - periodic ping-pong to continuously refresh neighbor lists
 - list size specified by user at peer: heterogeneity means some peers may have more neighbors
 - Gnutella found to follow power law distribution
 - $P(\#links = L) \sim L^{-k}$ (k is constant)
- problems
 - ping/pong constituted 50% traffic
 - solution: multiplex, cache and reduce frequency of pings/ pongs
 - multiplex make one pong message out of multiple songs or pings
 - repeated searches with same keywords
 - solution: cache query, query hit messages
 - modem-connected hosts do not have enough bandwidth for passing Gnutella traffic
 - solution: use a central server to act as proxy for such peers
 - another solution:
 - fast track system
- freeloaders
 - user who only downloads never uploads
- Flooding causes excessive traffic
 - is there some way of maintaining meta-information about peers that leads to more intelligent routing?
 - structured peer to peer systems

FastTrack and Bit torrent

FastTrack

- hybrid between Gnutella and Napster
- some peers designated as super nodes
- supernode
 - stores a directory listing a subset of nearby (<filename>, <peer pointer>), similar to Napster servers
- any peer can become a super node, provided it has earned enough reputation

BitTorrent

- seed has full file
- leecher has some blocks of the file
- file split into blocks

- download local rarest first block policy: prefer early download of blocks that are least replicated among neighbors
- tit for tat bandwidth usage: Provide blocks to neighbors that provided it the best download rates
- choking: limit number of neighbors to which concurrent uploads \leq a number i.e., the “best” neighbors
 - every else choked
 - periodically re-evaluate this set
 - optimistic unchoke: periodically, unchoke a random neighbor helps keep unchoked set fresh

Chord

Distributed Hash Table

- allows you to insert, lookup, and delete object keys
- a distributed hash table allows you to do the same in a distributed setting
- performance concerns
 - load balancing
 - fault tolerance
 - efficiency of lookups and inserts
 - locality (messages that are transmitted are transmitted relatively close)
- Napster, Gnutella, FastTrack are all DHT's so is chord

Chord

- intelligent choice of neighbors to reduce latency and message
- uses consistent hashing
 - SHA-1(ip_address.port) -> 160 bit string
 - Truncated to m bits
 - called peer id(number between 0 and $2^m - 1$)
 - Not unique but id conflicts very unlikely
 - Can then map peers to one of 2^m logical points on a circle

Ring or Peers

- 2^m logical points
- has nodes onto the peer id on the circle
- e.x., node16 -> node32
- $m = 7$ so 0 to 127
- every node knows its immediate clockwise successor (ring) N112 -> N16

Finger Tables

- $m = 7$
- 7 finger tables
- ith entry at peer with id n is first peer with id $\geq n + (2^i)(\text{mod } 2^m)$
- warps around using modulus

What about files

- filenames also mapped using same consistent hash function
 - SHA-1(filename) -> 160 bit string (key)
 - File is stored at first peer with id greater than or equal to its key (mod 2^m)

- File cnn.com/index.html that maps to key K42 is stored at first peer with id greater than 42
 - Note that we are considering a different file-sharing application here: cooperative web caching

Search

- at node n , send query for key k to largest successor/finger entry $\leq k$ if none exist, send query to $\text{successor}(n)$
- takes $O(\log(N))$ time
 - intuition: at each step, distance between query and peer-with-file reduces by a factor of 2
 - intuition: after $\log(N)$ forwarding, distance to key is at most $2^m / 2^{\log(n)} = 2^m / N$
 - Number of node identifiers in a range of $2^m / N$ is $O(\log(N))$ with high probability
 - so using successors in that range will be ok, using another $O(\log(N))$ hops

Failures In Chord

search under peer failures

- one solution: maintain r multiple successor entries
- in case of failures, use successor entries
- choosing $r = 2\log(n)$ suffices to maintain lookup correctness
- if node fails that contains the file
- solution store the file at one successor and one predecessor

Need to deal with dynamic changes

- peers fails
- new peers join
- peers leave
 - P2P systems have a high rate of churn (node join, leave and failure)
 - lower in managed clusters
 - common feature in all distributed systems, including wide-area, clusters, clouds, etc
- So, all the time, need to
 - need to update successors and fingers, and copy keys

New Peers joining

- server gives ip address of some peer in the system
- routes message to self using hard routing protocol
 - makes it way to successor
- updates to successors and predecessors
- copies peers from successor as the peers in the system
- copies its successor finger table it uses it as it's own finger table
- stabilization protocol runs in the background (followed by all nodes)
 - asks immediate neighbors predecessors and successors for their finger tables
- may need to copy some files/keys from the clockwise node to itself

Stabilization protocol

- concurrent peer joins, leaves, failures might cause loopiness of pointers, and failure of lookups
 - chord peers periodically run a stabilization algorithm that checks and updates pointers and keys
 - ensures non-loopiness of fingers, eventual success of lookups and $O(\log(n))$ lookups
 - each stabilization round at a peer involves a constant number of messages

- strong stability takes $O(N^2)$ stabilization rounds

Churn

- when nodes are constantly joining, leaving, failing
 - significant effect to consider: traces from the Overnet system show hourly peer turnover rates (churn) could be 25-100% of total number of nodes in system
 - leads to excessive (unnecessary) key copying (remember that keys are replicated)
 - stabilization algorithm may need to consume more bandwidth to keep up
 - main issue is that files are replicated, while it might be sufficient to replicate only meta information about files
 - alternatives
 - introduce a level of indirection
 - replicate metadata more

Virtual Nodes

- Hash can get non-uniform -> Bad load balancing
 - treat each node as multiple virtual nodes behaving independently
 - each joins the system
 - reduces variance and reduces load imbalance

Pastry

- assign ids to nodes
- leaf set - each node knows its successor(s) and predecessor(s)
- routing tables (based on prefix matching) - think of hypercube
- routing is based on prefix matching is therefore $\log(n)$
 - and hops are short

Pastry routing

- id 01110100101
- It maintains a neighbor peer with an id matching each of the following prefixes
- when it needs to route to a peer 01110111001 it starts by forwarding to a neighbor with the largest matching prefix
- for each prefix, among all potential neighbors with a matching prefix, the neighbor with the shortest round trip time is selected
- since shorter prefixes have many more candidates, the neighbors for shorter prefixes are likely to be closer than the neighbors for longer prefixes
- Thus, in the prefix routing, early hops are short and later hops are longer

Summary of Chord and Pastry

- Chord and Pastry protocols
 - more structured than Gnutella
 - Black box lookups algorithms
 - Churn handling can get complex
 - $O(\log(N))$ lookup hops may be high
 - can we reduce the number of hops

Kelips - A 1 Hop lookup DHT

- k affinity groups
 - $k \sim \sqrt{N}$
- each node based to a group
- Nodes neighbors
 - almost all other nodes in its own affinity group
 - one contact node per foreign affinity group
- peer knows about all the other peers in its group
 - also knows about one contact member from another group

Kelips files and metadata

- file can be stored at any node
- decouple file replication/location from file querying
- each filename hashed to a group
 - all nodes in the group replicate pointer information and the file (meta information)
 - affinity group does not store files
- lookup
 - find file affinity group
 - go to your contact for the file affinity group
 - failing that try another of your neighbors to find a contact
- lookup = 1 hop (or a few)
 - memory cost $O(\sqrt{N})$
 - 1.93 mb for 100k nodes. 10m files
 - fits in ram of most workstations/laptops

Kelips Soft state

- membership lists
 - gossip based membership
 - within each affinity group
 - and also across affinity groups
 - $O(\log(n))$ dissemination time
- file meta data
 - needs to be periodically refreshed from source code
 - times out (deletes the information about the file)
- range of tradeoffs available
 - memory vs lookup cost vs background bandwidth