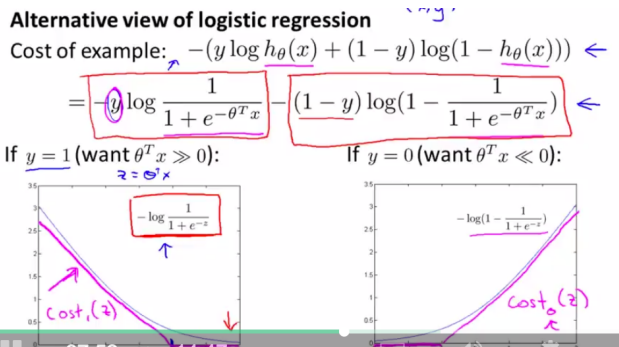


- Optimization objective
- Alternative view of logistic regression
  - sigmoid activation
  - if  $y = 1$ 
    - we want  $h_{\theta}(x)$  close to 1
    - $\theta^T x \gg 0$
  - if  $y = 0$ 
    - we want  $h_{\theta}(x)$  close to 0
    - $\theta^T x \ll 0$
- Alternative view of logistic regression
  - new cost function gives support vector machine advantages and easier optimization problem
  - two new cost functions one when  $y = 1$  and one when  $y = 0$



- Support vector machine
  - Logistic regression:

**Support vector machine**

Logistic regression:

$$\min_{\theta} \frac{1}{m} \left[ \sum_{i=1}^m y^{(i)} \underbrace{(-\log h_{\theta}(x^{(i)}))}_{\text{cost}_1(\theta^T x^{(i)})} + (1 - y^{(i)}) \underbrace{(-\log(1 - h_{\theta}(x^{(i)})))}_{\text{cost}_0(\theta^T x^{(i)})} \right] + \frac{\lambda}{2m} \sum_{j=1}^n \theta_j^2$$

Support vector machine:

$$\min_{\theta} C \sum_{i=1}^m y^{(i)} \text{cost}_1(\theta^T x^{(i)}) + (1 - y^{(i)}) \text{cost}_0(\theta^T x^{(i)}) + \frac{1}{2} \sum_{j=1}^n \theta_j^2$$

Handwritten notes and examples:

$\min_u (u-5)^2 + 10 \rightarrow u=5$

$\min_u 10(u-5)^2 + 10 \rightarrow u=5$

$A + \lambda B \leftarrow C$

$C = \frac{1}{\lambda}$

$\rightarrow \min_{\theta} C \sum_{i=1}^m [y^{(i)} \text{cost}_1(\theta^T x^{(i)}) + (1 - y^{(i)}) \text{cost}_0(\theta^T x^{(i)})] + \frac{1}{2} \sum_{j=1}^n \theta_j^2$

- replace the cost function with the new cost functions for support vector machines
- Logistic regression uses lambda to control the relative weighting between the first and second terms
  - large value of lambda gives B a very large weight
- Support vector machine uses c as an alternative parameter
  - we then minimize  $CA + B$ 
    - giving C a large value gives A a very large weight
- Different way of controlling or parameterizing the weights
  - $c = 1/\lambda$

- SVM hypothesis
  - does not output a probability
  - just makes a direct prediction if  $y = 1$  or  $y = 0$

### SVM hypothesis

$$\rightarrow \min_{\theta} C \sum_{i=1}^m \left[ y^{(i)} \text{cost}_1(\theta^T x^{(i)}) + (1 - y^{(i)}) \text{cost}_0(\theta^T x^{(i)}) \right] + \frac{1}{2} \sum_{j=1}^n \theta_j^2$$


Hypothesis:

$$h_{\theta}(x) = \begin{cases} 1 & \text{if } \theta^T x \geq 0 \\ 0 & \text{otherwise} \end{cases}$$

- Large Margin Intuition
- Support vector machine

### Support Vector Machine

$$\rightarrow \min_{\theta} C \sum_{i=1}^m \left[ y^{(i)} \text{cost}_1(\theta^T x^{(i)}) + (1 - y^{(i)}) \text{cost}_0(\theta^T x^{(i)}) \right] + \frac{1}{2} \sum_{j=1}^n \theta_j^2$$



$\rightarrow$  If  $y = 1$ , we want  $\theta^T x \geq 1$  (not just  $\geq 0$ )  $\theta^T x \geq 1$   
 $\rightarrow$  If  $y = 0$ , we want  $\theta^T x \leq -1$  (not just  $< 0$ )  $\theta^T x \leq -1$

- SVM Decision Boundary
  - set the constant  $C$  to be a very large value
  - if  $C$  is large highly motivated to choose a term to make sure the first term is equal to zero in the objective function
- SVM Decision Boundary: Linearly separable case
  - margin minimum distance to any of the training examples
  - large margin classifier
- Large Margin classifier in presence of outliers
  - if  $C$  is very large with change the decision the boundary for an outlier
  - if  $C$  is not very large the SVM will not change the decision boundary for an outlier
- Kernels
  - Given  $x$  compute new feature depending on proximity to landmarks  $l(1)$ ,  $l(2)$ , and  $l(3)$

Given  $x$ :  
 $f_1 = \text{similarity}(x, l^{(1)}) = \exp\left(-\frac{\|x - l^{(1)}\|^2}{2\sigma^2}\right)$   
 $f_2 = \text{similarity}(x, l^{(2)}) = \exp\left(-\frac{\|x - l^{(2)}\|^2}{2\sigma^2}\right)$   
 $f_3 = \text{similarity}(x, l^{(3)}) = \exp(\dots)$   
 Kernel (Gaussian kernels)

- Kernels and Similarity
  - measure how similar  $x$  is to one of the landmarks

#### Kernels and Similarity

$$f_1 = \text{similarity}(x, l^{(1)}) = \exp\left(-\frac{\|x - l^{(1)}\|^2}{2\sigma^2}\right) = \exp\left(-\frac{\sum_{j=1}^n (x_j - l_j^{(1)})^2}{2\sigma^2}\right)$$


If  $x \approx l^{(1)}$ :

$$f_1 \approx \exp\left(-\frac{0^2}{2\sigma^2}\right) \approx 1$$


If  $x$  is far from  $l^{(1)}$ :

$$f_1 = \exp\left(-\frac{(\text{large number})^2}{2\sigma^2}\right) \approx 0.$$

- each landmark can now compute new features given the corresponding landmark
- varying sigma
  - .5 the feature falls to zero more rapidly
  - 3 the value of the feature falls away much more slowly



- Kernels 2
  - landmarks allowed us to define the similarity function or the kernel
  - take the examples and for every training example we are just going to call it put landmarks as the same locations as the training examples
  - features are going to measure how close something is as seen in training set



- SVM with Kernels
  - Given  $m$  training examples
  - choose  $m$  landmarks

- Given example  $x$ :

Given example  $x$ :

$$\begin{aligned} \rightarrow f_1 &= \text{similarity}(x, l^{(1)}) \\ \rightarrow f_2 &= \text{similarity}(x, l^{(2)}) \\ &\vdots \end{aligned}$$

$$f = \begin{bmatrix} f_0 \\ f_1 \\ f_2 \\ \vdots \\ f_m \end{bmatrix} \quad f_0 = 1$$

- For training example

For training example  $(x^{(i)}, y^{(i)})$ :

$$x^{(i)} \rightarrow \begin{bmatrix} f_1^{(i)} \\ f_2^{(i)} \\ \vdots \\ f_m^{(i)} \end{bmatrix} \quad \begin{aligned} f_1^{(i)} &= \sin(x^{(i)}, l^{(1)}) \\ f_2^{(i)} &= \sin(x^{(i)}, l^{(2)}) \\ &\vdots \\ f_m^{(i)} &= \sin(x^{(i)}, l^{(m)}) \end{aligned}$$

$x^{(i)} \in \mathbb{R}^{n+1}$  (or  $\mathbb{R}^n$ )

$$f^{(i)} = \begin{bmatrix} f_0^{(i)} \\ f_1^{(i)} \\ f_2^{(i)} \\ \vdots \\ f_m^{(i)} \end{bmatrix} \quad f_0^{(i)} = 1$$

- SVM with Kernels
  - Hypothesis: Given  $x$ , compute features  $f$  element of  $\mathbb{R}^{m+1}$
  - Training:
    - to obtain the parameters for the support vector machine solve the minimization problem

Training:

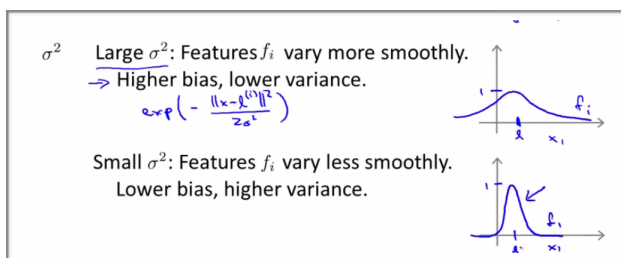
$$\min_{\theta} C \sum_{i=1}^m y^{(i)} \text{cost}_1(\theta^T f^{(i)}) + (1 - y^{(i)}) \text{cost}_0(\theta^T f^{(i)}) + \frac{1}{2} \sum_{j=1}^m \theta_j^2$$

$\theta = \begin{bmatrix} \theta_0 \\ \theta_1 \\ \vdots \\ \theta_m \end{bmatrix}$  (ignore  $\theta_0$ )

$M = 10,000$

$$\begin{aligned} - \sum_j \theta_j^2 &= \theta^T \theta \leftarrow \theta = \begin{bmatrix} \theta_1 \\ \vdots \\ \theta_m \end{bmatrix} \\ - &\rightarrow \theta^T M \theta \leftarrow \|\theta\|^2 \end{aligned}$$

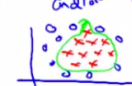
- SVM parameters
  - $C = 1$  lambda
  - Large  $C$ : Lower bias, high variance (small lambda)
  - Small  $C$ : Higher bias, low variance (large lambda)



- Using an SVM
- Support vector machine
  - choice of parameter C
  - choice of kernel
  - e.g. No kernel ("linear kernel")
    - predict "y = 1" if  $\theta^T x \geq 0$

E.g. No kernel ("linear kernel")  
 Predict "y = 1" if  $\theta^T x \geq 0$   $\theta_0 + \theta_1 x_1 + \dots + \theta_n x_n \geq 0$   $x \in \mathbb{R}^{n+1}$   
 $\rightarrow n \text{ large, } m \text{ small}$

Gaussian kernel:  
 $f_i = \exp\left(-\frac{\|x - l^{(i)}\|^2}{2\sigma^2}\right)$ , where  $l^{(i)} = x^{(i)}$ .  $x \in \mathbb{R}^n, n \text{ small}$   
 Need to choose  $\sigma^2$ .  $\text{and/or } m \text{ large}$



- Kernel (similarity) functions:

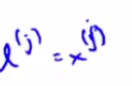
**Kernel (similarity) functions:**

function  $f = \text{kernel}(x_1, x_2)$

$f = \exp\left(-\frac{\|x_1 - x_2\|^2}{2\sigma^2}\right)$

return

$x \rightarrow \begin{matrix} f_1 \\ f_2 \\ \vdots \\ f_m \end{matrix}$



Note: Do perform feature scaling before using the Gaussian kernel.

- feature scaling is important
- Other choices of kernel
  - Not all similarity functions  $\text{similarity}(x, l)$  make valid kernels. (need to satisfy technical condition called "Mercer's Theorem" to make sure SVM packages' optimizations run correctly, and do not diverge).
  - polynomial kernel:  $k(x, l) = (x'l)^2, (x'l)^3, (x'l + 1)^3, (x'l + 5)^4$
  - more esoteric: String kernel, chi-square kernel, histogram intersection kernel
- Logistic regression VS SVMs

#### Logistic regression vs. SVMs

$n$  = number of features ( $x \in \mathbb{R}^{n+1}$ ),  $m$  = number of training examples

$\rightarrow$  If  $n$  is large (relative to  $m$ ): (e.g.  $n \geq m$ ,  $n = 10,000$ ,  $m = 10 - 1000$ )

$\rightarrow$  Use logistic regression, or SVM without a kernel ("linear kernel")

$\rightarrow$  If  $n$  is small,  $m$  is intermediate: ( $n = 1 - 1000$ ,  $m = 10 - 10,000$ )

$\rightarrow$  Use SVM with Gaussian kernel

If  $n$  is small,  $m$  is large: ( $n = 1 - 1000$ ,  $m = 50,000 +$ )

$\rightarrow$  Create/add more features, then use logistic regression or SVM without a kernel

$\rightarrow$  Neural network likely to work well for most of these settings, but may be slower to train.

