

Vector space model

- matching word more times more credit
- some words are more important than other words

Improved Vector Placement

- Term frequency vector
 - $\text{Sim}(q,d) = q \cdot d$ (dot product)
- inverse document frequency
 - count the number of documents that don't contain a particular word
- document frequency
 - count of documents that contain a particular word

IDF (inverse document frequency) (common words have low IDF and rare words have high IDF)

- $\log[(M+1)/k]$
 - k = total number of docs containing W (doc frequency)
 - M = total number of docs in collection

How effective is VSM with TF-IDF weighting

- fixing one problem tends to lead to another problem

Ranking function with TF-IDF Weighting

- $f(q,d) = \sum_{i=1}^N x_i y_i$
 - = summation all matched query words in document d
 - $c(w,q) c(w,d) (\log(M+1)/df(w))$

TF Transformation: $c(w,d) \rightarrow TF(w,d)$

- linear $y = x$
- 0/1 bit vector
- $y = \log(1+x)$
 - controls the inference of high weight
- $y = \log(1 + \log(1 + x))$
- BM25 transformation (best function)
 - $y = ((k+1)x)/(x+k)$
 - upper bounded by $k + 1$
 - varying k can simulate different transformation functions
 - setting $k = 0$ turns it into a bit vector 0/1 transformation
 - setting k to very large number makes the transformation look like a linear transformation
 - upper bound controls the inference from the upper bound
- Sub linear TF Transformation is need to
 - capture the intuition of diminishing return from higher TF
 - avoid dominance by one single term over all others
- Ranking function with BM25

— —

Doc Length Normalization

- penalize a long doc with a length normalizer
- pivoted length normalizer: average doc length as pivot
 - normalizer = 1 if $l_d = \text{average doc length}$

Document length normalizer

- if length of document is longer than the pivot then there is a penalty
- $\text{normalizer} = 1 - b + b \cdot (dl/avdl)$ b element of $[0, 1]$
- b controls the normalization
 - $b = 0$ no normalization
 - $b > 0$ the value would be higher for longer documents and reward for short documents.
Nothing if equal
- adjusting B we can control the degree of normalization

State of the Art VSM Ranking Functions

- pivoted length normalization VSM
 - summation of all matched query words in the documents $c(w, q)$
 - TF Transformation on the top. Double In In
 - document length normalizer on the bottom (using the pivot)
 - inverse document frequency normalizer
- BM25
 - inverse document frequency normalizer
 - summation of all matched query words in the documents $c(w, q)$
 - middle normalization
 - BM25 TF normalization
 - then a document length normalization multiplied by the BM 25 normalizer on the bottom

Further improvement

- improved instantiation of dimension?
 - stemmed words, stop word removal, phrases, latent semantic indexing, character n-grams
 - bag of words sufficient
 - language-specific and domain specific tokenization is important to ensure “normalization of terms”
- improved instantiation of similarity function?
 - cosine of angle between two vectors?
 - euclidean?
 - dot product seems still the best (sufficiently general especially with appropriate term weighting)

BM25

- further improvement?
 - combine the frequency counts of terms in all fields and then apply BM25
 - address the problem of over penalization of long documents by BM25 by adding a small constant to TF

Summary

- $\text{Relevance}(q, d) = \text{similarity}(q, d)$
- query and documents represent as vectors
- Heuristic design of ranking function.
- Major term weighting heuristics
 - TF, IDF, and document length normalization
- BM25 and Pivoted normalization seem to be most effective
 - BM25 derived using probabilistic modeling

Typical IR System Architecture

- three parts
 - indexer
 - scorer
 - feedback
 - online or offline
- tokenization
 - normalize lexical units = words with similar meanings should be mapped to the same indexing term
 - stemming: mapping all inflectional forms of words to the same root form
 - some languages pose challenges in word segmentation
- indexing
 - indexing = convert documents to data structures that enable fast search
 - inverted index is the dominating indexing method for supporting basic search algorithms
 - other indices may be needed for feedback
- inverted index
 - dictionary and postings
 - dictionary - contains the count of documents and the total frequency
 - postings - contains the document id and the frequency
 - can also store the positions of the words
- inverted index for fast search
 - single term query?
 - multi-term boolean query?
 - multi-term keyword query?

Empirical distribution of words

Zipf's law

- rank * frequency \sim constant
- $F(w) = C / r(w)^\alpha$

Data Structures for Inverted Index

- Dictionary: modest size
 - fast random access
 - preferred to be in memory
 - Hash table, B-tree, trie
- Postings: huge
 - sequential access is expected
 - can stay on disk
 - may contain docID, term frequency, term pos, etc
 - compression is desirable

Inverted Index Construction

- memory based methods: not usable for large collections
- sort-based methods:
 - collect local (termID, docID, freq) tuples
 - sort local tuples (to make "runs")
 - pair-wise merge runs

- output inverted file

Sort-based Inversion

- obtain all the information containing the
 - document identification
 - term identification
 - count
- merge sort so that the entries are sorted based on term identifications

Inverted Index Compression (Encodings)

- leverage skewed distribution of values and use variable length encoding
- TF compression
 - small numbers tend to occur far more frequently than large numbers
 - fewer bits for small (high frequency) integers at the cost of more bits for large integers
- Doc id compression
 - “d-gap” (store difference): $d_1, d_2-d_1, d_3-d_2, \dots$
 - feasible due to sequential access

Integer Compression Methods

- binary encoding
- unary encoding
- gamma encoding

Uncompress Inverted Index

- decoding of encoded integers
 - unary decoding
 - gamma decoding
- decode doc ids encoded using d-gap

— — —

How to Score Documents Quickly?

- general form of scoring function

$$f(q,d) = f_a(h(g(t_1,d,q), \dots), f_d(d), f_q(q))$$

- f_a
 - combines all the functions
- h
 - inside of function
 - functions that compute the weights of contribution of matched query term in d
 - inside functions weight aggregation
- f_d, f_q
 - adjustment/scoring factors of document and query

General algorithm for ranking documents

- $f_d(d)$ and $f_q(q)$ are precomputed
- maintain a score accumulator for each d to compute h
- for each query term t_i
 - fetch the inverted list $\{(d_1, f_1), \dots, (d_n, f_n)\}$
 - for each entry (d_j, f_j) , compute $g(t_i, d_j, q)$, and update score accumulator for doc d_j to incrementally compute h
- adjust the score to compute the function f_a then sort

Ranking based on TF Sum

- $f(d, q) = g(t_1, d, q) + \dots$ where $g(t_i, d, q) = c(t_i, d)$

Further Improving Efficiency

- Caching
- Keep only the most promising accumulators
- Scaling up to the Web-scale?