

## - **Distributed computing and apache spark**

- How to handle massive data
  - need more hardware to store/process modern data
  - scale-up (one big machine)
    - can be very fast for medium scale problems
    - expensive, specialized hardware
    - eventually hit a wall
  - scale-out (many small machines, i.e. distributed)
    - commodity hardware, scales to massive problems
    - need to deal with network communication
    - added software complexity
- What is apache spark
  - general, open-source cluster computing engine
  - well-suited for machine learning
    - fast iterative computations
    - efficient communication primitives
  - Simple and Expressive
    - APIs in Scala, Java, Python, R
    - Interactive shell
  - Integrated Higher-Level Libraries
    - Spark SQL, Spark streaming, MLib, GraphX
  -

## - **What is Machine Learning**

- A definition
  - Constructing and studying methods that learn from and make predictions on data
  - Broad area involving tools and ideas from various domains
    - Computer Science
    - Probability and Statistics
    - Optimization
    - Linear Algebra
- Some Examples
  - Face Recognition
  - Link Prediction
  - Text or document classification
  - Protein structure prediction
  - Games
- Terminology
  - Observations. Items or entities used for learning or evaluation e.g., emails
  - Features. Attributes (typically numeric) used to represent an observation, e.g. length, date, presence of keywords
  - Labels. Values/ categories assigned to observations e.g. spam, not-spam
  - Training and Test data. Observations used to train and evaluate a learning algorithm, e.g., a set of emails along with their labels
    - Training data is given to the algorithm for training
    - Test data is withheld at train time (use a validation set before test set)
- Two Common Learning Settings
  - Supervised Learning. Learning from labeled observations
    - Labels 'teach' algorithm to learn mapping from observations to labels
  - Unsupervised learning. Learning from unlabeled observations

- Learning algorithm must find latent structure from features alone
- Can be goal in itself (discover hidden patterns, exploratory data analysis)
- Can be means to an end (preprocessing for supervised task)
- Examples of Supervised Learning
  - Classification: Assign a category to each item, e.g. spam detection
    - Categories are discrete
    - Generally no notion of 'closeness' in multi-class setting
  - Regression. Predict a real value for each item, e.g. stock prices
    - Labels are continuous
    - Can define 'closeness' when comparing predictions with labels
- Examples of Unsupervised Learning
  - Clustering. Partition observations into homogeneous regions, e.g., to identify "communities" within large groups of people in social networks
  - Dimensionality Reduction. Transform an initial feature representation into a more concise representation, e.g., representing digital images
- **Typical Supervised Learning Pipeline**
  - Obtain Raw Data => Feature Extraction
    - Initial observations can be in arbitrary format
    - We extract features to represent observations
    - We can incorporate domain knowledge
    - We typically want numeric features
    - Success of entire pipeline depends on choosing good descriptions of observations
  - Feature Extraction <=> Unsupervised Learning
    - Preprocessing step for supervised learning
  - Obtain Raw Data => Feature Extraction => Supervised Learning
    - Train a supervised model using labeled data, e.g., Classification or Regression model
  - Obtain Raw Data => Feature Extraction => Supervised Learning => Evaluation
    - How do we determine the quality of the model we've just learned?
    - We can evaluate it on test / hold-out data, i.e., labeled data not used for training
    - If we don't like the results, we iterate
      - Feature Extraction => Supervised Learning => Evaluation
  - Obtain Raw Data => Feature Extraction => Supervised Learning => Evaluation => Predict
    - Once we're happy with our model, we can use it to make predictions on future observations, i.e., data without a known label
- **Sample Classification Pipeline**
  - Classification
    - Goal: Learn a mapping from observations to discrete labels given a set of training examples (supervised learning)
    - Example: Spam Classification
      - Observations are emails
      - Labels are {spam, not-spam} (Binary Classification)
      - Given a set of labeled emails, we want to predict whether a new email is spam or not-spam
  - Other examples
    - Fraud detection: User activity -> {Fraud, not fraud}
    - Face detection: Images -> set of people
    - Link prediction: Users -> {Suggest link, don't suggest link}
    - Clickthrough rate prediction: User and ads -> {click, no click}
  - Classification Pipeline

- Obtain Raw data
  - Raw data consists of a set of labeled training observations
- Feature extraction
  - typically transforms each observations into a vector of real numbers (features)
  - success or failure of a classifier often depends on choosing good descriptions of observations
  - “Bag of Words”
    - Observations are documents
    - Build vocabulary
    - Derive feature vectors from Vocabulary
- Supervised Learning: Train classifier using training data
  - Common classifiers include Logistic regression, SVMs, Decision trees, Random Forests, etc.
  - Training (especially at scale) often involves iterative computations, e.g., gradient descent
  - Logistic Regression
    - Goal: Find linear decision boundary
      - Parameters to learn are feature weights and offset
      - Nice probabilistic interpretation
      - Covered in more detail later in course
- Evaluation
  - How can we evaluate the quality of our classifier?
    - We want good predictions on unobserved data
      - ‘Generalization’ ability
    - Accuracy on training data is overly optimistic since classifier has already learned from it
      - we might be ‘overfitting’
  - Overfitting and Generalization
    - Fitting training data does not guarantee generalization, e.g., lookup table
    - Left: perfectly fits training samples, but it is complex / overfitting
    - Right: misclassifies a few points, but simple / generalizes
    - Occam’s razor
  - How can we evaluate the quality of our classifier?
    - Idea: Create test set to simulate unobserved data
  - Evaluation: Split dataset into training / testing datasets
    - Train on training set (don’t expose test set to classifier)
    - Make predictions on test set (ignoring test labels)
    - Compare test predictions with underlying test labels
    - Various metrics to compare predicted and true labels
    - Accuracy is common for classification
  - Predict: Final classifier can then be used to make predictions on future observations, e.g., new emails we receive
- Linear algebra review
  - Matrices
  - Vectors
    - a matrix with many rows and one column
  - Transpose
    - swap the rows and columns of the matrix
  - Addition and subtraction
    - Element wise operations

- matrices must have the same dimensions
- Matrix scalar multiplication
  - we multiply each matrix element by the scalar value
- Scalar Product
  - Function that maps two vectors to a scalar
  - Performs pairwise multiplication of vector elements
  - Take the sum
  - The two vectors must be the same dimension
  - Also known as dot product or inner product
- Matrix-Vector Multiplication
  - Matrix A and vector w
    - output y
  - $y_i$  equals the scalar product between the with row of A and w
  - we repeat for each row of A, so if A has n rows, does y
  - To perform inner products, # columns in A must equal # rows of w
- Scalar Product Revisited
  - Special case of Matrix-Vector Multiplication
    - Vectors assumed to be in column form (many rows, one column)
    - Transposed vectors are row vectors
    - Common notation for scalar product:  $x'w$
- Matrix Matrix Multiplication
  - Takes first row of first matrix scalar product of first column gives first row column for new matrix then second ...
  - Take the second row ...
  - To perform inner products the # columns in A must equal # rows of B
  - $m \times n * n \times p = m \times p$
- Outer Product
  - Special case of Matrix Matrix Multiplication involving two vectors
  - $C_{ij}$  is the product of it entry of x and the nth entry of w
- Identity Matrix
  - One is the scalar multiplication identity
  - Identity matrices are square, with ones on the diagonal entries
- Inverse Matrix
  - Multiplying a matrix by its inverse results in the identity matrix
  - only a square matrix can have an inverse
- Euclidean Norm for Vectors
  - The magnitude / length of a scalar is its absolute value
  - Vector norms generalize this idea for vectors
  - The Euclidean norm for x element  $R^m$  is denoted by  $\|x\|_2$

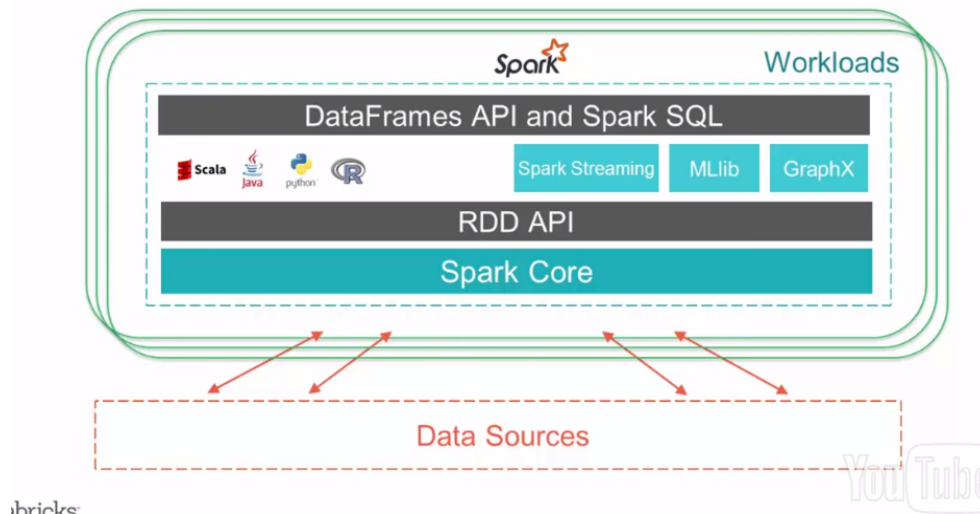
- $\|x\|_2 = \sqrt{x_1^2 + x_2^2 + \dots + x_m^2}$
- Equals absolute value when  $m=1$
- Related to scalar product:  $\|x\|_2^2 = x^T x$

## - Big O Notation for Time and Space Complexity


- Big O Notation
  - describes how algorithms respond to changes in input size
    - Both in terms of processing time and space requirements

- We refer to complexity and Big O notation synonymously
- Required space proportional to units of storage
  - typically eight bytes to store a floating point number
- Required time proportional to number of basic operations
  - arithmetic operations (+, -, x, /), logical tests (<, >, ==)
- Notation:  $f(x) = O(g(x))$ 
  - can describe an algorithm's time or space complexity
    - Informal definition: f does not grow faster than g
    - Formal definition:  $|f(x)| \leq C|g(x)|$  for all  $x > N$
    - Ignores constants and lower-order terms
      - For large enough x, these terms won't matter
- $O(1)$  Complexity
  - Constant time algorithms perform the same number of operations every time they're called
  - E.g. performing a fixed number of arithmetic operations
  - Similarly, constant space algorithms require a fixed amount of storage every time they're called
    - E.g. storing the results of a fixed number of arithmetic operations
- $O(n)$  Complexity
  - Linear time algorithms perform a number of operations proportional to the number of inputs
    - E.g. adding two n-dimensional vectors requires  $O(n)$  arithmetic operations
  - Similarly, linear space algorithms require storage proportional to the size of the inputs
    - E.g. adding two n-dimensional vectors results in a new n-dimensional vector which requires  $O(n)$  storage
- $O(n^2)$  Complexity
  - Quadratic time algorithms perform a number of operations proportional to the square of the number of inputs
    - E.g. outer product of two n-dimensional vectors requires  $O(n^2)$  multiplication operations (one per each entry of the resulting matrix)
  - Similarly, quadratic space algorithms require storage proportional to the square of the size of the inputs
    - E.g. outer product of two n-dimensional vectors requires  $O(n^2)$  storage (one per each entry of the resulting matrix)
- Time and Space Complexity Can Differ
  - Inner product of two n-dimensional vectors
    - $O(n)$  time complexity to multiply n pairs of numbers
    - $O(1)$  space complexity to store result (which is scalar)
  - Matrix inversion of an  $n \times n$  matrix
    - $O(n^3)$  time complexity to perform inversion
    - $O(n^2)$  space complexity to store result
- Matrix-Vector Multiply
  - Goal: multiply an  $n \times m$  matrix an  $m \times 1$  vector
  - Computing result takes  $O(nm)$  time
    - There are n entries in the resulting vector
    - each entry computed via dot product between two m-dimensional vectors (a row of input matrix and input vector)
  - Storing result takes  $O(n)$  space
    - The result is an n-dimensional vector

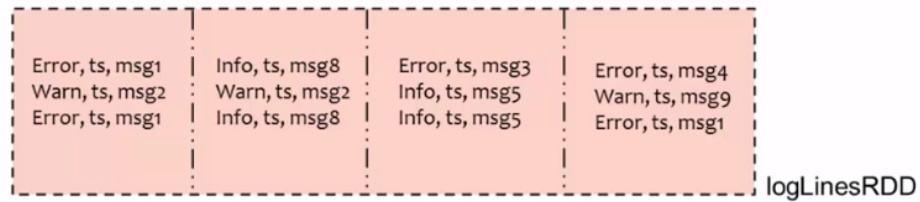
- E.g. Matrix-Matrix Multiply
  - Computing result takes  $O(npm)$  time
    - There are  $np$  entries in the resulting matrix
    - Each entry computed via dot product between two  $m$ -dimensional vectors
  - Storing takes  $O(np)$  total
- **RDD Fundamentals**



- Driver program to worker machines
- Resilient Distributed Datasets (RDDs)
  - Write programs in terms of operations on distributed data
  - Partitioned collections of objects spread across a cluster
  - Diverse set of parallel transformations and actions
  - Fault tolerant



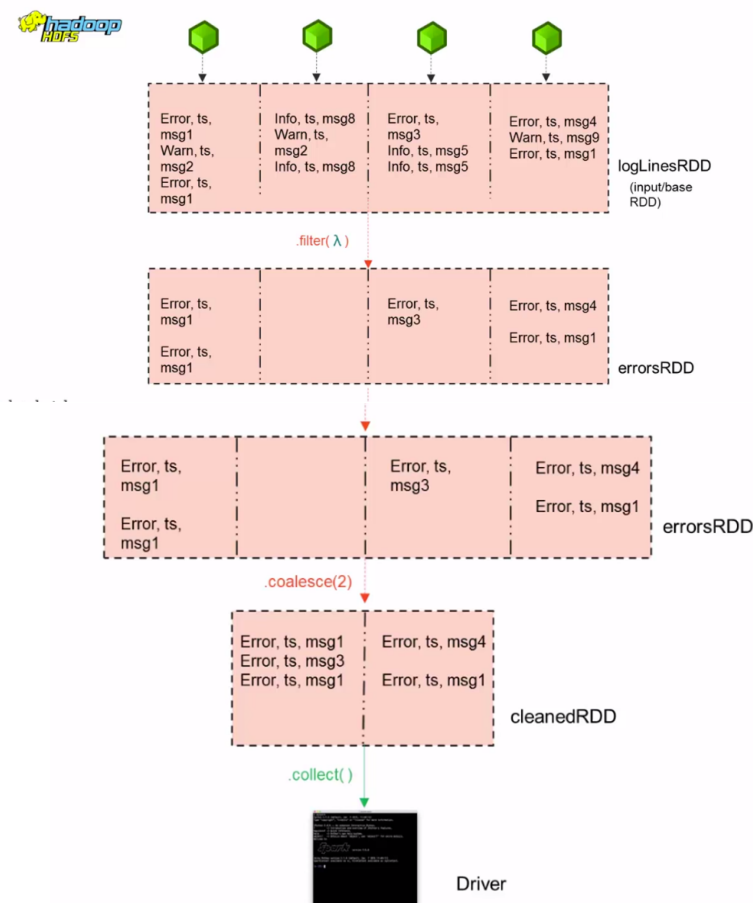
RDD w/ 4 partitions



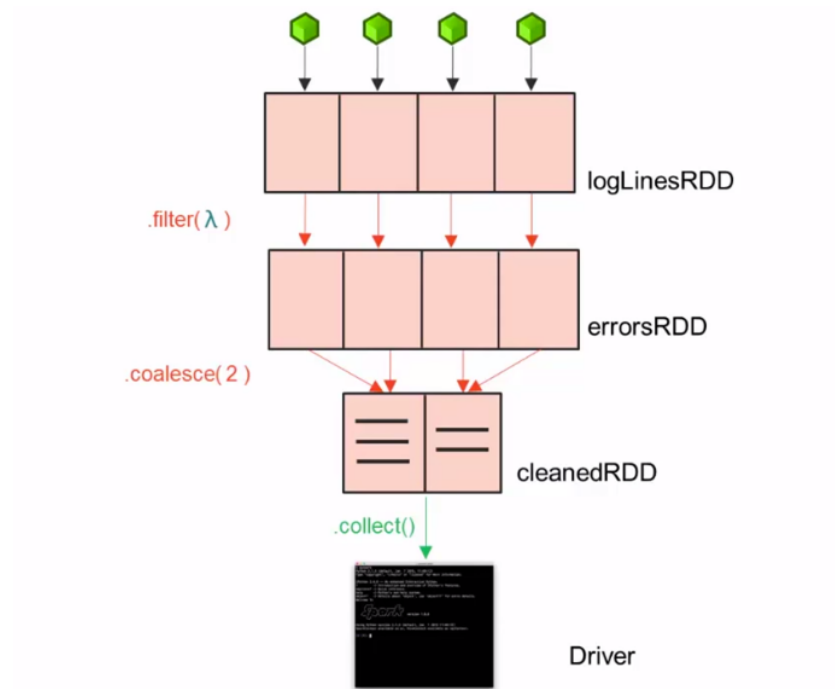
A base RDD can be created 2 ways:

- Parallelize a collection
- Read data from an external source (S3, C\*, HDFS, etc)


- Create a Base RDD
  - Parallelize in Python
    - wordsRDD = sc.parallelize(["fish", "cats", "dogs"])
    - Parallelize - Take an existing in-memory collection and pass it to SparkContext's parallelize method
  - Read a local txt file in python
    - linesRDD = sc.textFile("/path/to/README.md")
    - Read from Text File
    - There are other methods to read data from HDFS, C\*, S3, HBase, etc.
- Operations on Distributed Data
  - Two types of operations: transformations and actions
  - Transformations are lazy (not computed immediately)
  - Transformations are executed when an action is run
  - Persist (cache) distributed data in memory or disk



## - DAG



## - Execution




- Collect operation completed





- Partition >>> Task >>> Partition

## Partition >>> Task >>> Partition



- Lifecycle of an RDD-based Spark Program
  - Create base RDD
  - Chain together transformations
    - example filter map
  - Cache intermediate RDDs
  - Perform actions
    - count collect to start parallel computation
- Transformations - are lazy and will not be computed until an action is called
  - Map()
  - intersection()
  - zipWithIndex()
  - flatMap()
  - distinct()
  - pipe()
  - filter()
  - groupByKey()
  - coalesce()
  - mapPartitions()
  - reduceByKey()
- Actions - result in a spark job being launched and cause any related transformations to be computed
  - reduce()
  - collect()
  - count()
  - first()
  - take()
  - takeOrdered()
  - saveAsTextFile()
- RDDs vs DataFrames
  - RDDs provide a low-level interface into Apache Spark
  - DataFrames have a schema
  - DataFrames are cached using Tungsten format
  - DataFrames are optimized via Catalyst
  - DataFrames are built on top of the RDD and core APIs

-  
-  
-  
-

- 
- 
- 
- 
- 
- sdf