

- Key Data Management Concepts
 - A data model
 - is a collection of concepts for describing data
 - A schema
 - is a description of a particular collection of data using a given data model
- Structure Spectrum
 - Structured
 - schema first
 - relational database
 - Semi-structured
 - schema later
 - documents xml
 - Unstructured
 - schema never
 - plain text
- Semi-Structured Tabular dat
 - A table is a collection of rows and columns
 - Each column has a name
 - Each cell may or may not have a value
 - Each column has a type (String, integer)
 - together the column type are the schema for the data
 - Two choices for how the schema is determined
 - spark dynamically infers the schema while reading each row
 - programmer statically specifies the schema
- Structured Data
 - a relational data model is the most used data model
 - relation,, a table with rows and columns
 - every relation has a schema defining each columns' type
 - the programmer must statically specify the schema
- Unstructured data
 - only one column with string or binary type
- The structure spectrum
 - extract-transform-load
 - impose structure on unstructured data

Analysis, Big Data, and Apache Spark

- What is Apache Spark?
 - scalable, efficient analysis of big data
- Some Traditional Analysis Tools
 - Unix shell commands, pandas, R
 - all run on a single machine
- Real world spark analysis use cases
 - big data genomics using adam
 - data processing for wearables and internet of things
- The Big Data Problem
 - Data growing faster than computation speeds
 - Growing data sources
 - Storage getting cheaper
 - But, stalling CPU speeds and storage bottlenecks
 - One machine can not process or even store all the data

- solution is to distribute data over cluster of machines
- the key is it is all about memory
- Distributed Memory
 - Big data partition into multiple frames
 - sparks data frame
- The Spark Computing Framework
 - Provides a programming abstraction and parallel runtime to hid complexities of fault-tolerance and slow machines
- Apache Spark Components
 - Spark SQL
 - Spark Streaming
 - MLib and ML (Machine Learning)
 - GraphX (graph)
 - Apache Spark
- Python Spark (pySpark)
 - pySpark provides an easy-to-use programming abstraction and parallel runtime
 - “heres an operation, run it on all of the data”
 - DataFrames are the key concept
- Spark Driver and Workers
 - A Spark program is two programs
 - A driver program and a workers program
 - Worker programs run on cluster nodes or in local threads
 - DataFrames are distributed across workers
- Spark and SQL Contexts
 - A Spark program first creates a SparkContext object
 - SparkContext tells Spark how and where to access a cluster
 - pySpark shell, Databricks CE automatically create SparkContext
 - iPython and programs must create a new SparkContext
 - The program next creates a sqlContext object
 - use sqlContext to create DataFrames
- Spark Essentials: Matter
 - The master parameter for a SparkContext determines which type and size of cluster to use

Apache Spark DataFrames

- DataFrames
 - The primary abstraction in Spark
 - Immutable once constructed
 - Track lineage information to efficiently recompute lost data
 - Enable operations on collection of elements in parallel
 - You construct DataFrames
 - by parallelizing existing Python collections
 - by transforming an existing Spark or pandas DF
 - from files in HDFS or any other storage system
 - Each row of a DataFrame is a Row object
 - The fields of a Row can be accessed like attributes
 - Two types of operations: transformations and actions
 - Transformations
 - are lazy (not computed immediately)
 - transformed DF is executed when action runs on it
 - Persist (cache) DFs in memory or disk

- Working with DataFrames
 - create DataFrame from data source
 - apply transformations to DataFrame: select filter
 - Apply actions to DataFrame: show count
- Creating DataFrames
 - Create DataFrames from Python collections(lists)
- pandas: Python Data Analysis Library
 - Open source data analysis and modeling library
 - an alternative to using R
 - pandas DataFrame: a table with named columns
 - the most commonly used pandas object
 - represented as python Dict(column_name -> series)
 - Each pandas series Object represents a column
 - 1-D labeled array capable of holding any data type
 - R has a similar data frame type
- Creating DataFrames
 - easy to create pySpark DataFrames from pandas DataFrames
 - `spark_df = sqlContext.createDataFrame(pandas_df)`
 - From HDFS, text files, JSON files, Apache Parquet, Hypertable, Amazon S3, Apache Hbase, SequenceFiles, any other Hadoop InputFormat, and directory or glob wildcard;
 - `df = sqlContext.read.text("sdfs")`
 - `df.collect()`
- Creating a DataFrame from a file
 - `distFile = sqlContext.read.text("../")`
 - loads text file and returns a DataFrame with a single string column named "value"
 - each line in text file is a row
 - lazy evaluation means no execution happens now

Apache Spark Transformations

- Spark Transformations
 - Create a new DataFrame from an existing one or source
 - lazy evaluation result not computed right away
 - spark remembers set of transformations applied to base DataFrame
 - spark uses Catalyst to optimize the required calculations
 - spark recovers from failures and slow workers
 - Think of this as a recipe for creating result
- Column transformations
 - apply method creates a DataFrame from one column
 - `ageCol = people.age`
 - You can select one or more columns from a DataFrame
 - `df.select(**)`
 - * selects all the columns
 - `df.select('name','age')`
 - * selects the name and age columns
 - `df.select(df.name, (df.age + 10).alias('age'))`
 - * selects the name and age columns, increments the values in the age column by 10, and renames (alias) the age + 10 columns as age
- More Column transformations
 - The drop method returns a new DataFrame that drops the specified column:
 - `df.drop(df.age)`

- Review: Python lambda Functions
 - small anonymous functions (not bound to a name)
 - lambda a, b: a + b
 - returns the sum of its two arguments
 - can use lambda functions wherever function objects are required
 - restricted to a single expression
- User Defined Function Transformations
 - transform a DataFrame using User Defined Function
 - from pyspark.sql.types import IntegerType
 - slen = udf(lambda s: len(s), IntegerType())
 - df.select(slen(df.name).alias('slen'))
 - * creates a DataFrame of [Row(slen=5), Row(slen=3)]
 - UDF takes named or lambda function and the return type of the function
- Other Useful transformations
 - filter(func)
 - returns a new DataFrame formed by selecting those rows of the source on which function returns true
 - where(func)
 - where is an alias for filter
 - distinct()
 - return a new DataFrame that contains the distinct rows of the source DataFrame
 - orderBy(*cols, **kw)
 - returns a new DataFrame sorted by the specified column(s) and in the sort order specified by kw
 - sort(*cols, **kw)
 - Like orderBy.sort returns a new DataFrame sorted by the specified column(s) and in the sort order specified by kw
 - explode(col)
 - returns a new row for each element in the given array or map
 - func is a Python named function or lambda function
- Using transformations
 - df = sqlContext.createDataFrame(data, ['name', 'age'])
 - [Row(name=u'Alice', age=1), Row(name=u'Bob', age=2)]
 - from pyspark.sql.types import IntegerType
 - doubled = udf(lambda s: s * 2, IntegerType())
 - df2 = df.select(df.name, doubled(df.age).alias('age'))
 - [Row(name=u'Alice', age=2), Row(name=u'Bob', age=4)]
 - *selects the name and age columns, applies the UDF to age column and alias resulting column to age
 - df3 = df2.filter(df2.age > 3)
 - example
 - data2
 - df = sqlContext.createDataFrame(data2, ['name', 'age'])
 - df2 = df.distinct()
 - * only keeps rows that are distinct
 - df3 = df2.sort("age", ascending=False)
 - * sort ascending on the age column
 - example
 - data3 = [Row(a=1, intlist=[1,2,3])]

- `df4 = sqlContext.createDataFrame(data3)`
 - `[Row(a=1, intlist=[1,2,3])]`
- `df4.select(explode(df4.intlist).alias("anInt"))`
 - `[Row(anInt=1), Row(anInt=2), Row(anInt=3)]`
 - * turn each element of the int list column into a Row, alias the resulting column to `anInt`, and select only that column
- GroupedData Transformations
 - `groupBy(*cols)` groups the DataFrame using the specified columns, so we can run aggregation on them
 - `agg(*exprs)`
 - compute aggregates (avg, max, min, sum, or count) and returns the result as a DataFrame
 - `count()`
 - counts the number of records for each group
 - `avg(*args)`
 - computes average values for numeric columns for each group
- Using GroupedData (1)
 - Example
 - data
 - `df = sqlContext.createDataFrame(data, ['name', 'age', 'grade'])`
 - `df1 = df.groupBy(df.name)`
 - `df1.agg({"*": "count"}).collect()`
 - Example (same results as the previous example)
 - data
 - `df = sqlContext.createDataFrame(data, ['name', 'age', 'grade'])`
 - `df.groupBy(df.name).count()`
 - Example (computes the average across all the numerics)
 - data
 - `df = sqlContext.createDataFrame(data, ['name', 'age', 'grade'])`
 - `df.groupBy().avg().collect()`
 - `[Row(avg(age)=2.5, avg(grade)=7.5)]`
 - Example (computes the average for age and grade grouping by name)
 - data
 - `df = sqlContext.createDataFrame(data, ['name', 'age', 'grade'])`
 - `df.groupBy('name').avg('age', 'grade').collect()`
- Transforming a DataFrame
 - `linesDF = sqlContext.read.text('...')`
 - creates a lines DataFrame
 - each line of the file
 - `commentsDF = linesDF.filter(isComment)`
 - Lazy evaluation - means nothing executes - Spark saves recipe for transforming source

Apache Spark Actions

- Spark Actions
 - Cause Spark to execute recipe to transform source
 - Mechanism for getting results out of Spark
- Some useful actions
 - `show(n, truncate)`
 - prints the first n rows of the DataFrame
 - `take(n)`

- returns the first n rows as a list of Row
- collect() (should never use collection in production applications) (make sure will fit in driver program)
 - return all the records as a list of Row
- count()
 - returns the number of rows in this DataFrame
 - count for DataFrames is an action, while for GroupedData it is a transformation
- describe(*cols)
 - Exploratory Data Analysis function that computes statistics (count, mean, stddev, min, max) for numeric columns - if no columns are given, this function computes statistics for all numerical columns
- Getting Data Out of DataFrames
 - Example
 - data
 - df = sqlContext.createDataFrame(data, ['name','age'])
 - df.collect()
 - [Row(name=u'Alice', age=1), Row(name=u'Bob', age = 2)]
 - df.show()
 - df.count()
 - df.take(1)
 - df.describe() (only works on numerical columns)
- Spark Programming Model
 - linesDf = sqlContext.read.text('....')
 - print linesDF.count()
 - count() causes Spark to:
 - read data
 - sum within partitions
 - combine sums in driver
 - commentsDF = linesDF.filter(isComment)
 - print linesDF.count()
 - commentsDF.count()
 - Spark recomputes linesDF:
 - read data (again)
 - sum within partitions
 - combine sums in driver
- Spark Program Lifecycle
 - Create DataFrames from external data or createDataFrame from a collection in driver program
 - Lazily transform them into new DataFrames
 - cache() some DataFrames for reuse
 - Perform actions to execute parallel computation and produce results

Best Programming Practices

- Local or Distributed?
 - Where does code run?
 - Locally, in the driver
 - Distributed at the executors
 - Both at the driver and the executors
- Very important question:
 - Executors run in parallel

- Executors have much more memory
- Where Code Runs
 - Most Python code runs in driver
 - except for code passed to transformations
 - Transformations run at executors
 - Actions run at executors and driver
- Examples
 - `a = a + 1`
 - runs in Your application (driver program)
 - `linesDF.filter(isComment)`
 - Spark executor
 - `commentsDF.count()`
 - Runs at both driver and executor
- How Not to Write Code
 - Let's say you want to combine two DataFrames: `aDF`, `bDF`
 - You remember that `df.collect()` returns a list of Row, and in Python you can combine two lists with `+`
 - A naive implementation would be
 - `a = aDF.collect()`
 - `b = bDF.collect()`
 - `cDF = sqlContext.createDataFrame(a+b)`
 - Where does this code run?
- `a + b`
 - `a = aDF.collect()`
 - `b = bDF.collect()`
 - all distributed data for `a` and `b` is sent to driver
 - What if `a` and/or `b` is very large?
 - Driver could run out of memory:
 - Out Of Memory error(OOM)
 - Also, takes a long time to send the data to the driver
- `a + b`
 - `cDF = sqlContext.createDataFrame(a+b)`
 - all data for `cDF` is sent to the executors
 - What if the list `a + b` is very large?
 - Driver could run out of memory:
 - Out of Memory error(OOM)
 - Also, takes a long time to send the data to executors
- The Best Implementation
 - `cDF = aDF.unionAll(bDF)`
 - use the DataFrame reference API
 - `unionAll()`:
 - return a new DataFrame containing union of rows in this frame and another frame
 - Runs completely at executors:
 - Very scalable and efficient
- Some Programming Best Practices
 - Use Spark Transformations and Actions wherever possible
 - Never use `collect()` in production, instead use `take(n)`
 - `cache()` DataFrames that you reuse a lot
-
